

# Project 1A

## COM S 362

### Fall 2021

**Note:** Iteration 1A must be completed by yourself. Only future project iterations will allow working in pairs (teams of 2).

## 1. Introduction

For this project, you will be creating your own system call in xv6. The new system call, `pcount`, will return the total number processes currently on the system.

## 2. Setup

Follow the steps in the Xv6 Development Setup guide linked to by Module 2 in Canvas. You can use an alternative setup, for example, installing xv6 on your local machine.

## 3. Organization of Xv6

The xv6 source is divided into two folders: `kernel` and `user`. As their names imply `kernel` is the part of the operating system that executes in kernel mode and `user` contains several utility programs including a simple shell.

To implement a system call we need to bridge the gap between user and kernel. Recall, a simple function call will not work because a user mode application does not have the privileges required to jump to instructions in the kernel's code. A system call is invoked when the user mode application executes a privileged instruction, which cause a trap, resulting in one of the kernel's interrupt handlers being executed.

On the user side, we can see the calls to `ecall` (a privileged instructions) in the file `user/usys.S` for each of the system calls (e.g., `fork()`, `exit()`, `wait()` ...). We will not need to directly modify this assembly language file, it is automatically generated by a Perl script.

On the kernel side, the interrupt handler that is invoked on a user caused trap is `uservec` in `kernel/trampoline.S`. We can see that it saves the user's CPU registers (context), calls `usertrap()` located in the file `kernel/trap.c`, and then restores the user's CPU registers.

The function `usertrap()` calls `syscall()` in `kernel/syscall.c` which uses the system call number to call the correct function from the array of function pointers `syscalls[]`.

## 4. Adding a New System Call (Kernel Side)

Now you will add and test your own system call in xv6! Start by modifying the following files.

### **kernel/syscall.h**

Add a new system call number using the name `SYS_pcount`. Follow the example of the other system calls in `syscall.h`. The system call number is used to look up which function to call from the `syscalls[]` array of function pointers, we will modify that next.

### **kernel/syscall.c**

Declare the kernel side function that will be called: `sys_pcount(void)`. Search for another system call function such as `sys_uptime` and follow the same pattern for `sys_pcount`. You will need to add code to two places in this file.

### **kernel/proc.h**

The Process Control Blocks (PCBs) in xv6 are stored in an array called `proc` (we will refer to this as the process table). To get access to the process table from any file in the kernel, add the following line to the bottom of `kernel/proc.h`:

```
extern struct proc proc[NPROC];
```

### **kernel/sysproc.c**

Now create the full definition for the `sys_count` function in `sysproc.c`. Again, search for `sys_uptime` and use it as an example.

The purpose of `pcount` is to return the total number of processes on the system. All processes are maintained in the array of `proc` structs called `proc[]`. The size of the array is `NPROC`. Iterate through the array to find all array entries that have a state that is not `UNUSED`. The state of `UNUSED` indicates the array entry is empty, it does not store a process.

## 5. Testing the System Call (User Side)

### **user/usys.pl**

Add the new system call to the Perl script that automatically generates the required assembly code. Follow the example of `uptime` to add an entry for `pcount`.

**user/user.h**

Search for `uptime` and add a similar function declaration for `pcount`.

## **Makefile**

Now we will make a new utility program to test the system call. In the `Makefile` search for `zombie` and add a similar line for `pctest`.

**user/pctest.c**

Copy the code from `user/zombie.c` to a new file called `user/pctest.c`.

Add statements to print the process count before calling `fork` and after the parent calls `sleep` have the parent print the process count. Are both counts the same or different, what should be the expected result and why? Add a comment to the code answering this question.

## 6. Submission

Document `user/pctest.c` with your name and answer to the question above, no other documentation is required. Submit a zip file of the `xv6-riscv` directory. On pyrite the zip file can be created using:

```
$ zip -r xv6-riscv.zip xv6-riscv
```

Submit `xv6-riscv.zip`.